Semantic Segmentation using FCNs

Computer Vision

Assignment 4

Varun Edachali (2022101029)

January 2025

1 Dataset Visualisation

We read the segmentation mask image, extract the first channel, isolate each of the 13 classes individually by creading binary masks, and visualise each binary mask.

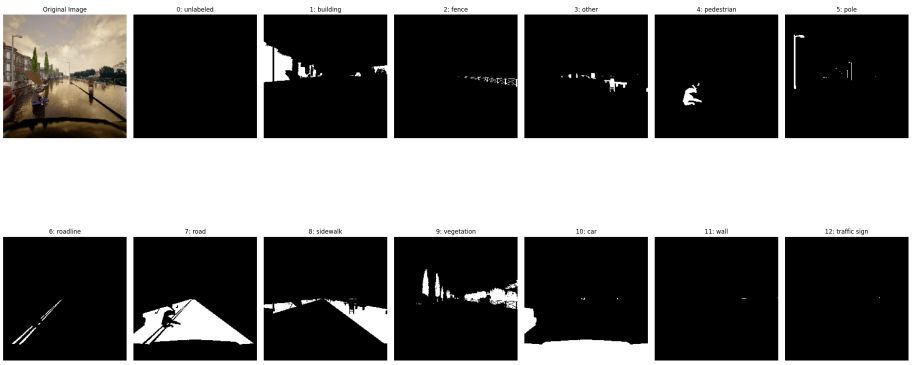

Figure 1: eda for fcn

2 FCN Variants

2.1 Frozen VGG Backbone

On freezing the VGG backbone, we see the following results:

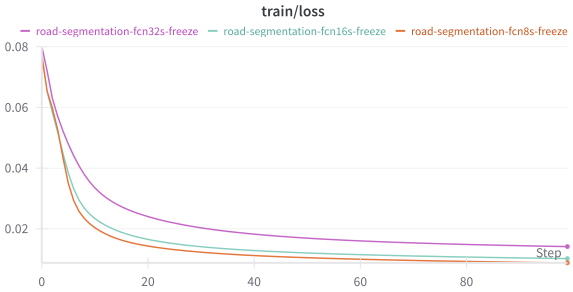

Figure 2: train loss - frozen backbone

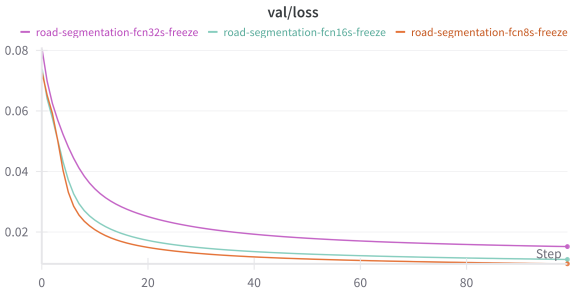

Figure 3: validation loss - frozen backbone

Figure 4: test loss - frozen backbone

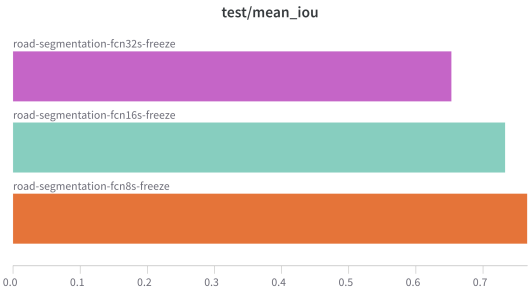

Figure 5: test mIoU - frozen backbone

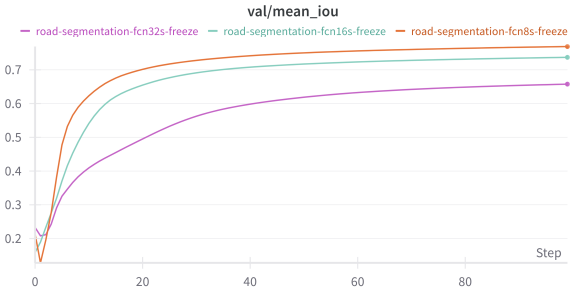

Figure 6: val mIoU - frozen backbone

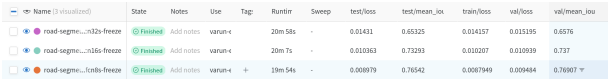

Figure 7: compiled results - frozen backbone

In addition, we visualise these predictions on 8 images from the test set, one of which is shown here. These images are available in full resolution in the repository in the `ckpts` dir for both questions.

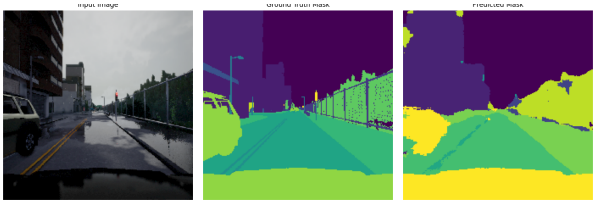

Figure 8: FCN8s - frozen backbone

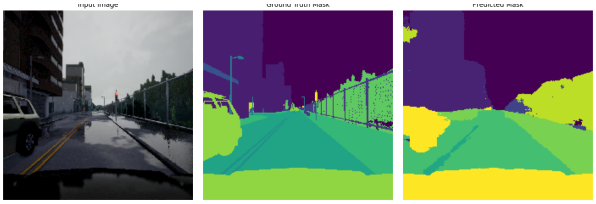

Figure 9: FCN16s - frozen backbone

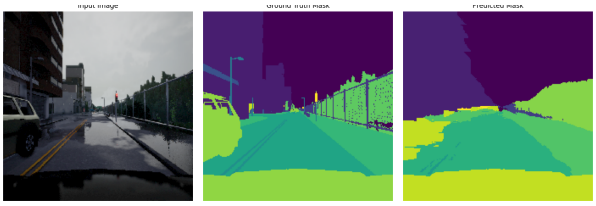

Figure 10: FCN32s - frozen backbone

2.2 UnFrozen VGG Backbone

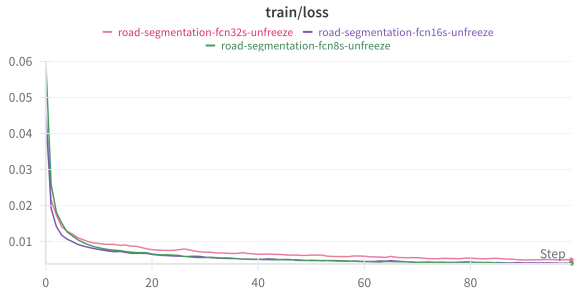

Figure 11: train loss - unfrozen backbone

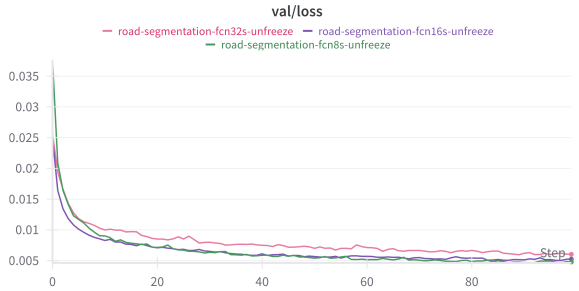

Figure 12: validation loss - unfrozen backbone

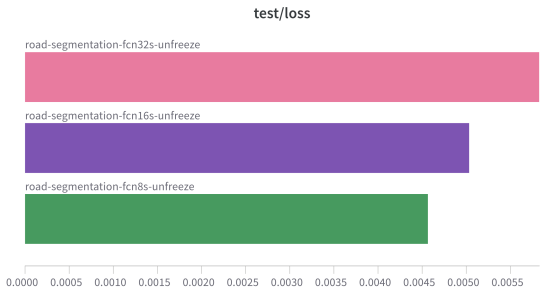

Figure 13: test loss - unfrozen backbone

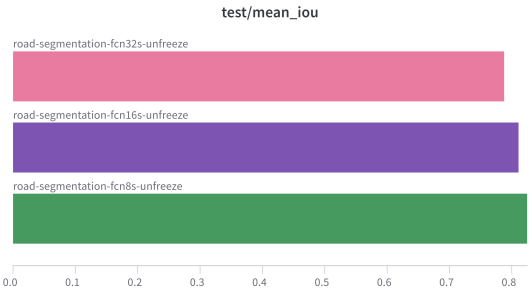

Figure 14: test mIoU - unfrozen backbone

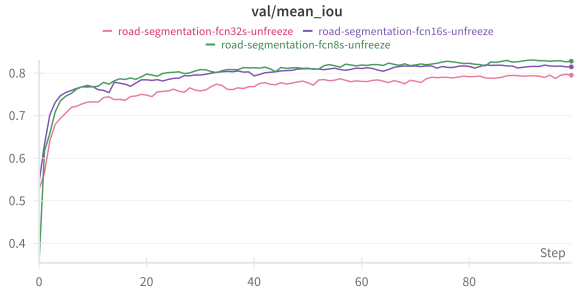

Figure 15: val mIoU - unfrozen backbone

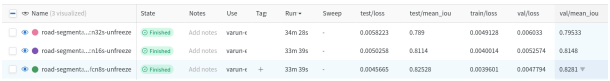

Figure 16: compiled results - unfrozen backbone

In addition, we visualise these predictions on 8 images from the test set, one of which is shown here. These images are available in full resolution in the repository in the `ckpts` dir for both parts.

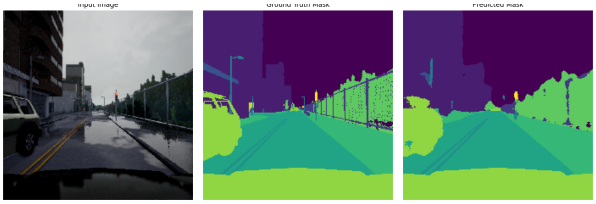

Figure 17: FCN8s - unfrozen backbone

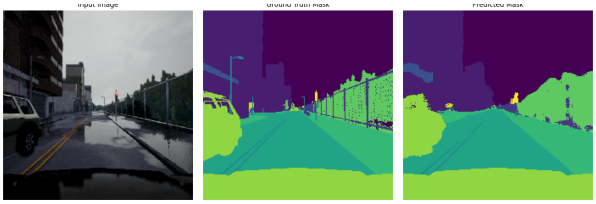

Figure 18: FCN16s - unfrozen backbone

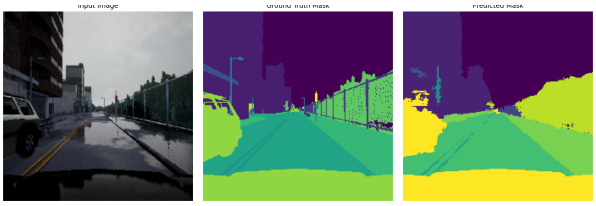

Figure 19: FCN32s - unfrozen backbone

2.3 Comparison

The core difference lies int he granularity of predictions / segmentations.

- **FCN-32s** converts the final fully connected layers into convolutional onces and directly upsamples the output in one step. This is simplest and fastest but produces course segmentations with less detail.
- **FCN-16s** improves upon FCN-32s by adding a skip connection from a shallower intermediate layer. This combination injects finer spatial details, leading to a better segmentation accuracy.
- **FCN-8s** further refines the segmentation by incorporating an additional skip connection from an even shallower layer. The fusion of features produces the finest spatial detail among the variants, leading to the best accuracy.

In addition, the rather obvious difference is in the role of the models. **FCN-8s** upsamples the features from $\frac{1}{8}$ resolution ($8\times$) to recover pixel level resolution, and so on.

We consolidate the results from above into a table to focus on the test mean IoU.

Table 1: Test Mean IoU for FCN Variants (Frozen vs. Unfrozen)

	fcn8s		**fcn16s**		**fcn32s**	
	Frozen	**Unfrozen**	**Frozen**	**Unfrozen**	**Frozen**	**Unfrozen**
test/mIoU	0.76542	0.82528	0.73293	0.8114	0.65325	0.789

We notice that the general order of performance is FCN-8s > FCN-16s > FCN-32s which aligns with our expectations. In addition, unfrozen > frozen in general, likely meaning the extracted VGG features do not sufficiently encode enough information for adequate segmentation, leading to the model learning better ones.

This aligns with the qualitative results as seen in the images in the above document and the repository. Qualitatively, the frozen models seem to make worse predictions overall - wrongly selected items being the cause of such issues more often than incorrect labelling, though. This could suggest some required features are not being encoded appropriately by the model, so the convolutionised classifier is unable to focus on them.